

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df=pd.read_csv("CarPrice_Assignment.csv")
```

```
In [3]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                205 non-null    int64
1   symboling              205 non-null    int64
2   CarName               205 non-null    object
3   fueltype              205 non-null    object
4   aspiration             205 non-null    object
5   doornumber            205 non-null    object
6   carbody               205 non-null    object
7   drivewheel            205 non-null    object
8   enginelocation        205 non-null    object
9   wheelbase             205 non-null    float64
10  carlength              205 non-null    float64
11  carwidth              205 non-null    float64
12  carheight              205 non-null    float64
13  curbweight             205 non-null    int64
14  enginetype            205 non-null    object
15  cylindernumber        205 non-null    object
16  enginesize             205 non-null    int64
17  fuelsystem            205 non-null    object
18  boreratio             205 non-null    float64
19  stroke                205 non-null    float64
20  compressionratio      205 non-null    float64
21  horsepower            205 non-null    int64
22  peakrpm               205 non-null    int64
23  citympg               205 non-null    int64
24  highwaympg            205 non-null    int64
25  price                 205 non-null    float64
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB
```

```
In [4]: df.corr()
```

Out[4]:

	car_ID	symboling	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	bore
car_ID	1.000000	-0.151621	0.129729	0.170636	0.052387	0.255960	0.071962	-0.033930	0.260064
symboling	-0.151621	1.000000	-0.531954	-0.357612	-0.232919	-0.541038	-0.227691	-0.105790	-0.130051
wheelbase	0.129729	-0.531954	1.000000	0.874587	0.795144	0.589435	0.776386	0.569329	0.488750
carlength	0.170636	-0.357612	0.874587	1.000000	0.841118	0.491029	0.877728	0.683360	0.606454
carwidth	0.052387	-0.232919	0.795144	0.841118	1.000000	0.279210	0.867032	0.735433	0.559150
carheight	0.255960	-0.541038	0.589435	0.491029	0.279210	1.000000	0.295572	0.067149	0.171071
curbweight	0.071962	-0.227691	0.776386	0.877728	0.867032	0.295572	1.000000	0.850594	0.648480
enginesize	-0.033930	-0.105790	0.569329	0.683360	0.735433	0.067149	0.850594	1.000000	0.583774
boreratio	0.260064	-0.130051	0.488750	0.606454	0.559150	0.171071	0.648480	0.583774	1.000000

stroke	-0.160824	-0.008735	0.160959	0.129533	0.182942	-0.055307	0.168790	0.203129	-0.05
compressionratio	0.150276	-0.178515	0.249786	0.158414	0.181129	0.261214	0.151362	0.028971	0.00
horsepower	-0.015006	0.070873	0.353294	0.552623	0.640732	-0.108802	0.750739	0.809769	0.57
peakrpm	-0.203789	0.273606	-0.360469	-0.287242	-0.220012	-0.320411	-0.266243	-0.244660	-0.25
citympg	0.015940	-0.035823	-0.470414	-0.670909	-0.642704	-0.048640	-0.757414	-0.653658	-0.58
highwaympg	0.011255	0.034606	-0.544082	-0.704662	-0.677218	-0.107358	-0.797465	-0.677470	-0.58
price	-0.109093	-0.079978	0.577816	0.682920	0.759325	0.119336	0.835305	0.874145	0.55

```
In [5]: df.drop(["drivewheel", "stroke", "doornumber", "carbody", "car_ID", "wheelbase", "cylindernumb
```

```
In [6]: df.drop(["CarName", "aspiration", "engine.location", "carheight"], axis=1, inplace=True)
```


```
In [7]: df
```

```
Out[7]:
```

	carlength	carwidth	curbweight	enginetype	enginesize	horsepower	peakrpm	citympg	highwaympg	
0	168.8	64.1	2548	dohc	130	111	5000	21	27	134
1	168.8	64.1	2548	dohc	130	111	5000	21	27	165
2	171.2	65.5	2823	ohcv	152	154	5000	19	26	165
3	176.6	66.2	2337	ohc	109	102	5500	24	30	135
4	176.6	66.4	2824	ohc	136	115	5500	18	22	174
...
200	188.8	68.9	2952	ohc	141	114	5400	23	28	168
201	188.8	68.8	3049	ohc	141	160	5300	19	25	190
202	188.8	68.9	3012	ohcv	173	134	5500	18	23	214
203	188.8	68.9	3217	ohc	145	106	4800	26	27	224
204	188.8	68.9	3062	ohc	141	114	5400	19	25	228

205 rows × 10 columns

```
In [8]: sns.boxplot(data=df);
```



```
In [9]: from sklearn.preprocessing import OrdinalEncoder
ordi=OrdinalEncoder()
```

```
df["enginetype"] = ordi.fit_transform(df[["enginetype"]])
```

```
In [10]: df.corr()
```

```
Out[10]:
```


	carlength	carwidth	curbweight	enginetype	enginesize	horsepower	peakrpm	citympg	highwaympg	price
carlength	1.000000	0.841118	0.877728	-0.113291	0.683360	0.552623	-0.287242	-0.670909	-0.704662	0.682920
carwidth	0.841118	1.000000	0.867032	0.012298	0.735433	0.640732	-0.220012	-0.642704	-0.677218	0.759325
curbweight	0.877728	0.867032	1.000000	-0.055265	0.850594	0.750739	-0.266243	-0.757414	-0.797465	0.835305
enginetype	-0.113291	0.012298	-0.055265	1.000000	0.040766	0.010301	0.005599	-0.085004	-0.078456	0.049171
enginesize	0.683360	0.735433	0.850594	0.040766	1.000000	0.809769	-0.244660	-0.653658	-0.677470	0.874145
horsepower	0.552623	0.640732	0.750739	0.010301	0.809769	1.000000	0.131073	-0.801456	-0.770544	0.808139
peakrpm	-0.287242	-0.220012	-0.266243	0.005599	-0.244660	0.131073	1.000000	-0.113544	-0.054275	-0.085267
citympg	-0.670909	-0.642704	-0.757414	-0.085004	-0.653658	-0.801456	-0.113544	1.000000	0.971337	-0.685751
highwaympg	-0.704662	-0.677218	-0.797465	-0.078456	-0.677470	-0.770544	-0.054275	0.971337	1.000000	-0.685751
price	0.682920	0.759325	0.835305	0.049171	0.874145	0.808139	-0.085267	-0.685751	-0.685751	1.000000

```
In [11]: df.drop(["enginetype", "peakrpm"], axis=1, inplace=True)
```

```
In [13]: q1=np.quantile(df["price"],0.25)
q3=np.quantile(df["price"],0.75)
iqr=q3-q1
uw=q3+1.5*iqr
lw=q1-1.5*iqr
for i in df["price"]:
    if i>uw:
        df["price"]=df["price"].replace(i,uw)
```

```
In [14]: sns.boxplot(data=df)
```

```
Out[14]: <AxesSubplot:>
```



```
In [15]: x=df.iloc[:, :-1]
y=df.iloc[:, -1]
```

```
In [16]: from sklearn.model_selection import train_test_split
xtrain,xtest,ytrain,ytest=train_test_split(x,y,test_size=0.3,random_state=1)
```

```
In [17]: from sklearn.linear_model import LinearRegression
```

```
lr=LinearRegression()  
lr.fit(xtrain,ytrain)  
ypred=lr.predict(xtest)
```

```
In [18]: from sklearn.metrics import mean_squared_error  
mse=mean_squared_error(ytest,ypred)  
rmse=np.sqrt(mse)  
print(rmse)
```

```
2673.985276526457
```

```
In [19]: lr.score(xtrain,ytrain)
```

```
Out[19]: 0.8271168025783082
```

```
In [20]: lr.score(xtest,ytest)
```

```
Out[20]: 0.8283586405637202
```

```
In [ ]:
```